

This lecture notes is based on Chapter 7 in “Algorithm Design” by Kleinberg and Tardos.


1. Flow networks

One often uses graphs to model transportation networks: networks whose edges carry some sort of traffic and whose nodes act as “switches” passing traffic between different edges, for example, a highway system in which edges are highways, and nodes are interchanges (cities), or a computer network in which edges are links carrying packets and nodes are switches.

Definition 1. A *flow network* is a directed graph $G=(V, E)$ with the following features:

- 1) associated with each edge e is a *capacity*, a non-negative number denoted as c_e ;
- 2) there is a single source node $s \in V$;
- 3) there is a single sink node $t \in V$.

The other nodes than s and t are called *internal nodes*.



Definition 2. In a flow network G , a *flow* is a function f that maps each edge e to a non-negative real number, $f: E \rightarrow \mathbb{R}^+$; the value $f(e)$ intuitively represents the amount of flow carried by edge e . A flow must satisfy the following properties:

- 1) (Capacity condition) for each $e \in E$, we have $0 \leq f(e) \leq c_e$;
- 2) (Conservation condition) for each node v other than s and t , we have

$$\sum_{e \in \text{In}(v)} f(e) = \sum_{e \in \text{Out}(v)} f(e), \text{ where } \text{In}(v) \text{ and } \text{Out}(v) \text{ represent the set of incoming and outgoing edges incident to } v, \text{ respectively.}$$


The value of a flow f , denoted by $v(f)$, is defined as the amount flow generated at s ,

which is equal to the flow entering t : $v(f) = \sum_{e \in \text{Out}(s)} f(e) = \sum_{e \in \text{In}(t)} f(e)$; this is because

$$\sum_{e \in \text{Out}(s)} f(e) = \sum_{v \in V} \sum_{e \in \text{Out}(v)} f(e).$$

Maximum-flow problem. Given a flow network, find a flow of maximum possible value.

Design an algorithm: we would like to push as much as flow from s to t .



Definition 3. Given a flow network G and a flow f , a *residual network (or graph)* G_f with respect to f is defined as,

- 1) The set of vertices is the same as G ;
- 2) For each edge $e=(u,v)$ in G on which $f(e)<c_e$, there are $c_e - f(e)$ “leftover” capacity; these edges are called the *forward edges*;
- 3) For each edge $e=(u,v)$ in G on which $f(e)>0$, there are $f(e)$ units of flow that we can “undo” (push backward) if we want to, thus we add an edge $e'=(v,u)$ in G_f with a capacity of $f(e)$; these edges are called the *backward edges*;

Let P be a simple s - t path (i.e. augmenting path) in G_f . We define $\text{bottleneck}(P, f)$ as the minimum residue capacity of any edge on P , w. r. t. flow f . Below we define an algorithm to find a new flow f' w.r.t. a given flow f and a path P .

- 1) let $b = \text{bottleneck}(P, f)$;
- 2) for each edge $(u, v) \in P$, if e is a forward edge, increase $f(e)$ in G by b ; otherwise decrease $f(e')$ by b , where $e'=(v,u)$;

Lemma 1. f' is a flow in G .

Proof: 1) each edge e in f' has capacity $\leq c_e$ because if e is a forward edge, $f'(e)=f(e)+b \leq f(e) + c_e - f(e) = c_e$; otherwise, $f'(e)=f(e)-b \geq 0$. 2). We only need to check the conservation condition on the internal nodes in P . Let v be such a node, the changes in the amount of flow entering v is the same as the changes of the amount of flow exiting v , depending on the types of incoming edge and outgoing edge (forward or backward).

Also because $v(f')=v(f)+\text{bottleneck}(P, f)$, and $\text{bottleneck}(P, f) > 0$, $v(f')>v(f)$. That means we always improve the flow by using the following algorithm.

Ford-Fulkerson algorithm for Maximum flow problem

Initially set $f(e)=0$ for all e in G

while there is a s - t path in G in the residue graph G_f

let P be a simple path from s to t in the residue graph G_f
 $f' \leftarrow \text{augment}(f, P)$
 $f \leftarrow f'$
 $G_f \leftarrow G_{f'}$
 Return f ;

Let all capacities in G are integers, and the flow out from s be $C = f^{\text{out}}(s) = \sum_{e \text{ out of } s} c_e$, which is the maximum value of any flow. Then Ford-Fulkerson algorithm terminates in at most C iterations. So the run time of Ford-Fulkerson algorithm is $O(mC)$.

Definition 4: In a flow network $G=(V, E)$, a *s-t cut* (A, B) is a partition of V into A and B so that $s \in A$ and $t \in B$. The capacity of a cut (A, B) , $c(A, B) = f^{\text{out}}(A) = \sum_{e \text{ out of } A} c_e$. The cut with the minimum capacity is called the *minimum cut* of G .

Lemma 2. Let f be any s-t flow, and (A, B) be any s-t cut. Then $v(f) = f^{\text{out}}(A) - f^{\text{in}}(A) = f^{\text{in}}(B) - f^{\text{out}}(B)$.

Proof: For any internal node v , $f^{\text{out}}(v) - f^{\text{in}}(v) = 0$; for s , $v(f) = f^{\text{out}}(s)$; $f^{\text{in}}(s) = 0$.

Corollary 1. Let f be any s-t flow, and (A, B) be any s-t cut. Then $v(f) \leq c(A, B)$.

Proof: $v(f) = f^{\text{out}}(A) - f^{\text{in}}(A) \leq f^{\text{out}}(A) = \sum_{e \text{ out of } A} c_e = c(A, B)$.

Corollary 2. If f is a maximum flow in G , there exists a s-t cut (A, B) in G for which $v(f) = c(A, B)$, and (A, B) is a minimum cut.

Proof: Let A be the set of nodes reachable by s in G_f , and $B = V - A$. All edges from a node in A to a node in B have capacity of c_e , and all edges from a node in B to a node in A has a capacity of 0. $v(f) = f^{\text{out}}(A) - f^{\text{in}}(A) = \sum_{e \text{ out of } A} c_e = c(A, B)$, i.e., $c(A, B)$ reaches the upper bound $v(f)$.




Ford-Fulkerson algorithm can also be used to find a minimum cut in a flow network.

Maximum flow minimum cut theorem. In every flow network, there is a flow f and a cut (A, B) so that $v(f) = c(A, B)$.

Note this theorem is not dependent on Ford-Fulkerson algorithm, and thus does not require the capacities to be integers.

2. Choosing good augmenting paths.

F-F algorithm run time: $O(mC)$. This could be very bad when C is a large integer.



Recall that augmentation increases the value of maximum flow by the bottleneck capacity of the augmenting path; so if we choose augmenting paths with large bottleneck capacity, we will make better progress. However, to find the path with the largest bottleneck capacity is hard. We will select a path with relatively large capacity, i.e., we maintain a scaling parameter Δ , and look for paths that have bottleneck capacity of at least Δ .

Let $G_f(\Delta)$ be the subgraph of residual graph G_f with capacity $\geq \Delta$: any augmenting path in this subgraph has bottleneck capacity $\geq \Delta$. We will work with Δ =powers of 2, and decrease it by half if we cannot find an augmenting path in $G_f(\Delta)$.

Scaling Max-flow Algorithm

Initially set $f(e)=0$ for all e in G

Set Δ to be the largest power of 2 that is smaller than C .

while $\Delta \geq 1$

 while there is a s-t path in G in the residue graph $G_f(\Delta)$

 let P be a simple path from s to t in the residue graph $G_f(\Delta)$

$f' \leftarrow \text{augment}(f, P)$

$f \leftarrow f'$

$G_f(\Delta) \leftarrow G_f(\Delta)$

$\Delta \leftarrow \Delta/2$

Return f ;

Observations: 1) the number of iterations of the outer loop is at most $1 + \lfloor \log_2 C \rfloor$. 2)

Within each inner loop, the flow value increases by at least Δ .

Lemma 3. Let f be the flow at the end of each inner loop. There is an s-t cut (A, B) in G for which $c(A, B) \leq v(f) + m\Delta$, and the maximum flow in the network has a value at most $v(f) + m\Delta$.

Corollary 3. The number of augmentations in the scaling phase is at most $2m$.

Proof: in last scaling phase, $\Delta' = 2\Delta$. Let f_p be the flow at the end of the last phase, the maximum value flow f^* is $v(f^*) \leq v(f_p) + m\Delta' = v(f_p) + 2m\Delta$. In the current scaling phase, each augmentation will increase the value of flow by Δ , so there are at most $2m$ augmentations.

Theorem. The run time of the scaling max-flow algorithm is at most $O(m^2 \log_2 C)$. When $C \gg m$, this is much faster than the max-flow algorithm $O(mC)$.

3. Preflow-push maximum-flow algorithms

Definition 4. A *preflow* in a flow network is a function f that maps each edge e to a non-negative real number, $f: E \rightarrow \mathbb{R}^+$, satisfying 1) the capacity conditions: for each e for each $e \in E$, we have $0 \leq f(e) \leq c_e$; and 2) for each node v other than the source s , we have

$\sum_{e \in \text{In}(v)} f(e) \geq \sum_{e \in \text{Out}(v)} f(e)$. We call the difference $e_f(v) = \sum_{e \in \text{In}(v)} f(e) - \sum_{e \in \text{Out}(v)} f(e)$ the *excess* of the preflow at v .

We can define the value of a preflow following definition 2, and a residual graph w.r.t. a preflow f , following definition 3.

Definition 5. A *labeling* (or *height*) is a function $h: V \rightarrow \mathbb{Z}^+$, from the nodes to non-negative integers. A labeling h and a preflow f are *compatible* if 1) $h(t)=0$, $h(s)=n$; and 2) (steepness condition) for all edges (v, w) in the residual graph, we have $h(v) \leq h(w)+1$.

Lemma 4. If a preflow f is compatible with a labeling h , there is no s - t path in residual graph G_f .

Proof: We prove it by contradiction. Let P be a simple s - t path in G_f , denoted by $s, v_1, \dots, v_k=t$. By definition, we have $h(s)=n$, $h(v_1) \geq h(s)-1=n-1$; \dots , $h(v_i) \geq n-i$. So $h(t) \geq n-k=0$. So $k=n$, i.e., the path visit one node more than one time, which is not a simple path. Contradiction.

Corollary 4. If s - t flow is compatible with a labeling h , f is a flow of maximum value.

Preflow-Push algorithm

Initiation: $h(s)=n$, for all other v , $h(v)=0$; $f(e)=c_e$ for all edges leaving s $e=(s,v)$ (these edges will not be in the residual graph), and $f(e)=0$ for all other edges.

Note: initial f and h are compatible.

Pushing and Relabeling: we want to turn preflow into a feasible flow, while keeping it compatible.

push(f, h, v, w) //pushing the excess of v , $e_f(v)$ along any edge (v,w) in the residual graph, where $h(w) < h(v)$, i.e. $h(v)=h(w)+1$.

if $e=(v,w)$ is a forward edge

increase $f(e)$ to $\min(c_e, f(e)+e_f(v))$

else if (v,w) is a backward edge

$e \leftarrow (w,v)$

decrease $f(e)$ to $\max(0, f(e)-e_f(v))$

relabel(f, h, v) //if cannot push the excess of v along any edge (v,w) in the residual graph, i.e., $e_f(v) > 0$ & for all edges (v,w) , $h(w) \geq h(v)$, we need to increase the label (height) of v . Increase $h(v)$ by 1

Preflow-Push

$h(s) \leftarrow n$, for all other v , $h(v) \leftarrow 0$;

$f(e) \leftarrow c_e$ for all edges leaving s , $e=(s,v)$, and $f(e)=0$ for all other edges.

while there is a node $v \neq t$ with $e_f(v) > 0$

 if there is w , s.t. $(v,w) \in E_f$, and $h(w) < h(v)$

 push(f, h, v, w)

 else

 relabel(f, h, v)

Observations: Throughout the algorithm, 1) the labels are non-negative integers; 2) f is a preflow with integer value; 3) the f and h are compatible; 4) the algorithm terminate when no node other than s or t has excess, so the final preflow f is flow.

Theorem. Preflow-Push algorithm terminates with a maximum flow.

Run time analysis.

Lemma 5. Let f be a preflow. If node v has excess, there is a path in G_f from v to the source s .

Proof: Let A be the set of nodes w such that there is a path from w to s in G_f and $B = V - A$. We need to show all nodes in B with excess of 0.

1) $s \in A$; and no edge e leaving A with $f(e) > 0$, otherwise there is a backward edge (y,x) in G_f and then y should be in A ;

2) Now consider the sum of excesses in B , recall that each node in B has nonnegative excess since s is not in B , so

$$0 \leq \sum_{v \in B} e_f(v) = \sum_{v \in B} (f^{in}(v) - f^{out}(v)) = -f^{out}(B), \text{ indicating the sum of excesses of nodes in } B = 0.$$

Lemma 6. Throughout the algorithm, all nodes have $h(v) \leq 2n-1$.

Proof: $h(t)=0$ and $h(s)=n$ do not change throughout the algorithm. For any other node v , its height can be increased by 1 only when it has excess (with preflow f), and then according to Lemma 5, there exist a path P from v to s in G_f . Let $|P|$ be the number of edges in the path. Along each edge (v,w) in the path, the height of nodes can decrease at most 1; hence, $h(v)-h(s) \leq |P| \leq n-1$.

Corollary 5. Throughout the algorithm, each node is relabeled at most $2n-1$ times, and the total number of relabeling operations is less than $2n^2$.

Definition 6. A push(f, h, v, w) operation is *saturating* if either $e=(v,w)$ is a forward edge in G_f and $f(e)$ increases c_e , or e is a backward edge and to 0, i.e., after the push, edge e is no longer in the residual graph G_f . All other push operations are called *nonsaturating*.

Lemma 7. Throughout the algorithm, the number of saturating push operations is at most $2nm$.

Proof: When we have a saturating push on (v, w) , we have $h(v)=h(w)+1$, and after the push, the edge (v, w) is no longer in G_f . To perform another push on (v, w) , we must make the edge appear again in G_f , i.e., to push from w to v , which requires to increase w 's height by 2 (so that w 's height above v 's). Notice that w 's height can increase at most $n-1$ times (Corollary 5); hence a saturating push from v to w can occur at most n times. Since every edge $e \in E$ can give rise to two edges (one forward and one backward) in the residual graph, overall, the saturating push can be performed at most $2mn$.


Lemma 8. Throughout the algorithm, the number of non-saturating push operations is at most $4n^2m$.

Proof: For a preflow f and a compatible labeling h , we define $\Phi(f, h) = \sum_{v: e_f(v) > 0} h(v)$ to be

the sum of heights of all nodes with positive excess. In the initial preflow and labeling, $\Phi(f, h)=0$; it remains nonnegative afterwards. When the algorithm terminates, $\Phi(f, h)=0$ again. A nonsaturating push decreases $\Phi(f, h)$ by at least 1, since after the push the node v will have no excess (get out of the sum) and $h(v)=h(w)+1 \geq 1$. Each relabeling will increase $\Phi(f, h)$ exactly by 1; so the total increment of $\Phi(f, h)$ by relabeling is at most $2n^2$. A saturating push may increase $\Phi(f, h)$ by the height of w , which is at most $2n-1$. There are at most $2nm$ saturation pushes, so the total increment of $\Phi(f, h)$ by saturating pushes is at most $2mn(2n-1)$. So $\Phi(f, h)$ can increase at most $4mn^2$ indicating there can be at most $4mn^2$ nonsaturating pushes.

Theorem. The run time of preflow-push algorithm is $O(mn^2)$.

4. bipartite matching problem and disjoint path problem.




Theorem. The size of the maximum matching in G is equal to the value of the maximum flow in G' ; and the edges in such a matching in G are the edges that carry flow from s to t

in G' .


F-F algorithm can solve maximum matching problem in $O(mn)$.

Definition 7. A set of paths in a graph G is *edge-disjoint* if their edges are disjoint, i.e., no two paths share an edge, though they may go through the same vertex.

Given a directed graph G with two distinct nodes s and t , the *directed edge-disjoint path problem* is to find the maximum number edge-disjoint s - t paths in G . Similarly, the *undirected edge-disjoint path problem* is to find the maximum number of edge-disjoint s - t paths in an undirected graph G .



Each directed graph G with two distinct nodes s and t can be converted into a flow network G' by assigning the capacity 1 to each edge in G , and setting s as the source and t as the sink.



Theorem. There are k edge-disjoint paths from s to t in a directed graph G from s to t iff the value of the maximum flow in G' is at least k .

Proof: \Rightarrow is straightforward, set $f(e)=1$ if e is included in one of the k edge-disjoint paths, and set $f(e)=0$ for remaining edges; This form a flow of value k .

\Leftarrow by induction on the value of flow k . if $k=0$, trivial. Otherwise there must be an edge (s,u) going out of s . We trace a path of edges that must also carry flow—there is always some edges carrying flow along the path—until either 1) we reach t ; or 2) we reach a vertex v that has been visited in the path. If 1), we find a path P from s to t . Let f' be the flow obtained by decreasing $f(e)$ for e in P . f' has a value at least $k-1$, according to induction, we have $k-1$ edge-disjoint paths in G using edges $\notin P$, and thus in total we have k edge-disjoint paths in G . If 2), eliminate the cycle from v to v and continue traverse the graph along the flow until reaching t . This will also give a edge-disjoint path.

Corollary 5. F-F algorithm can be used to find a maximum set of edge disjoint s - t paths in a directed graph G in $O(mn)$ time.

We can convert an undirected graph G into a directed graph G' by replacing each edge (u,v) in G with two directed edge (u,v) and (v,u) .

Lemma 9. In any flow network, there is maximum flow f where for all opposite directed edges $e=(u,v)$ and $e'=(v,u)$, either $f(e)=0$ or $f(e')=0$.

Proof: Assume f is a maximum flow. We can convert f to another flow f' that satisfying the condition: for every pair of e and e' , let $\delta=\min(f(e), f(e'))$, decrease both $f(e)$ and $f(e')$ by δ . It is easy to show that f' is a flow with the same value as f .


Corollary 6. There are k edge-disjoint paths in an undirected graph G from s to t iff the maximum value of an s - t flow in the directed version G' of G is at least k .

Corollary 7. F-F algorithm can be used to find a maximum set of edge disjoint s - t paths in an undirected graph G in $O(mn)$ time.

Network connectivity problem. Given a directed graph $G=(V,E)$ and two distinct nodes s and t in the graph, find a minimum number of edges whose removal will disconnect s and t .

Theorem (Menger, 1927) The maximum number of edge-disjoint path is equal to the minimum number of edges whose removal will disconnect s and t .

Proof: \Leftarrow , Assume F is the minimum set of edges whose removal disconnect s and t . Every path from s to t will use one edge e in F . So there are at most $|F|$ edge-disjoint paths. \Rightarrow Assume there are k edge-disjoint paths from s to t . Then the maximum flow in the graph G' is k , and thus the minimum cut (A, B) of the graph has capacity of k . This indicates there are k edges from A to B , which forms the set of edges whose removal will disconnect s and t .




5. Circulation with supplies and demands

Definition 8. A directed graph $G=(V,E)$ with *supplies and demands* is defined as a directed graph, with 1) edge capacities $c(e)$, for each $e \in E$, and 2) node supply and demands $d(v)$, $v \in V$: if $d(v)=0$ we call the node transshipment; if $d(v) > 0$; we call it demand; if $d(v) < 0$; we call it supply.

Definition 9. A *circulation* is a function that satisfies: 1) for each $e \in E$: $0 \leq f(e) \leq c(e)$

(capacity condition) and 2) for each $v \in V$: $\sum_{v \in V} f^{in}(v) - \sum_{v \in V} f^{out}(v) = d(v)$ (conservation condition)


Circulation problem: given (V, E, c, d) , is there a feasible circulation?



Lemma 10. (Necessary condition) if there is a feasible circulation with demands $\{d(v)\}$, then $\sum_v d(v) = 0$.

Convert circulation problem into a maximum flow problem. Given $G=(V,E)$, c and d , we build a flow network G' by

- 1) adding new source s and sink t ;
- 2) for each v with $d(v) < 0$, adding edge (s, v) with capacity $-d(v)$;
- 3) for each v with $d(v) > 0$, adding edge (v, t) with capacity $d(v)$.



Observation: the resulting graph is a flow network.

Lemma 11 Let D be the total demands in the graph: $D = \sum_{v:d(v)>0} d(v) = - \sum_{v:d(v)<0} d(v)$. The

graph G has a feasible circulation with demands $\{d(v)\}$ iff the maximum s - t flow in G' has value D (saturated).

Definition 10. A *circulation with lower bound* is a function that satisfies: 1) for each $e \in E$: $l(e) \leq f(e) \leq c(e)$ (capacity condition) and 2) for each $v \in V$:

$$\sum_{v \in V} f^{in}(v) - \sum_{v \in V} f^{out}(v) = d(v) \text{ (conservation condition).}$$

Circulation with lower bound problem: given (V, E, c, l, d) , is there a feasible circulation?

Convert circulation with lower bound problem into a circulation problem (i.e., modeling lower bound by demands).



6. Applications

Survey design

Design survey asking n_1 consumers about n_2 products.

Can only survey consumer i about product j if they own it.

Ask consumer i between c_i and c_i' questions, each for a different product.

Ask between p_j and p_j' consumers about product j .

Goal: design a survey that meets these specifics, if possible.

Special case: when $c_i = c_i' = p_i = p_i' = 1$, the problem can be formulated as a bipartite perfect matching.

Formulated as a circulation problem with lower bounds: include an edge (i, j) if consumer j owns product i , a circulation \leftrightarrow feasible survey design.

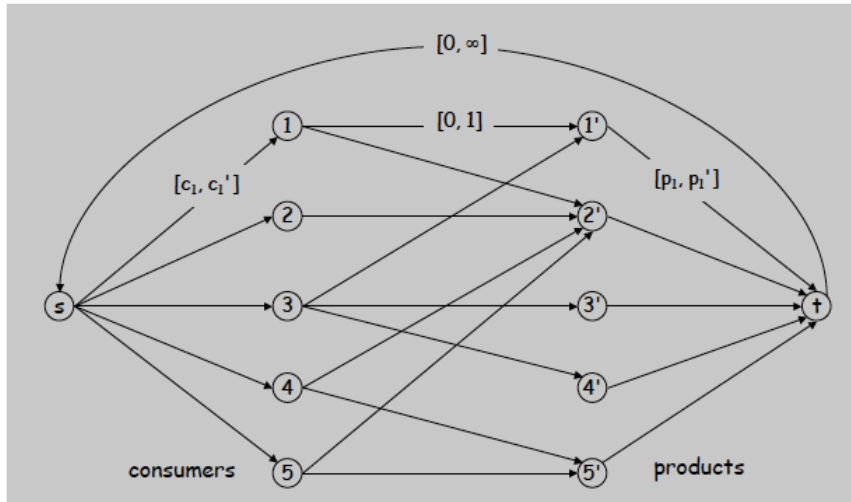


Image segmentation

To divide image into coherent regions, e.g., three people standing in front of complex background scene, to identify each person as a coherent object.

Problem formulation: foreground / background segmentation.

- 1) label each pixel in picture as belonging to foreground or background.
- 2) V = set of pixels, E = pairs of neighboring pixels.
- 3) $a_i \geq 0$ is likelihood pixel i in foreground; $b_i \geq 0$ is likelihood pixel i in background.
- 4) $p_{ij} \geq 0$ is separation penalty for labeling one of two neighboring pixels i and j as foreground, and the other as background.

Goals:

- 1) Accuracy: if $a_i > b_i$ in isolation, prefer to label i in foreground.
- 2) Smoothness: if many neighbors of i are labeled foreground, we should be inclined to label i as foreground.


Scoring: to find partition (A =foreground, B =background) that maximizes:


$$\sum_{i \in A} a_i + \sum_{i \in B} b_i - \sum_{\substack{(i,j) \in E, \\ |A \cap \{i,j\}| = 1}} p_{ij}, \text{ is equivalent to minimize } \sum_{i \in B} a_i + \sum_{i \in A} b_i + \sum_{\substack{(i,j) \in E, \\ |A \cap \{i,j\}| = 1}} p_{ij}.$$

Formulating it as a minimum cut problem:

- 1) $G' = (V', E')$.
- 2) add source to correspond to foreground;
- 3) add sink to correspond to background;
- 4) use two anti-parallel edges instead of undirected edge.

Then the minimum cut (A, B) have the minimum score as indicated above.






Project selection

We are given a set of projects with prerequisites

- 1) Set P of possible projects. Project v has associated revenue $p(v)$: some projects generate money: create interactive e-commerce interface, redesign web page; others cost money: upgrade computers, get site license
- 2) Set of prerequisites E . If $(v, w) \in E$, can't do project v and unless also do project w .
- 3) A subset of projects $A \subseteq P$ is feasible if the prerequisite of every project in A also belongs to A .


Goal: choose a feasible subset of projects to maximize revenue.

Build a prerequisite graph: include an edge from v to w if can't do v without also doing w , then $\{v, w, x\}$ is feasible subset of projects; $\{v, x\}$ is infeasible subset of projects.



Minimum cut formulation: constructing the flow network of G'

- 1) Assign capacity ∞ to all prerequisite edge.
- 2) Add two nodes s, t .
- 2) Add edge (s, v) with capacity $p(v)$ if $p(v) > 0$.
- 3) Add edge (v, t) with capacity $-p(v)$ if $p(v) < 0$.



Theorem: The cut (A, B) in G' is minimum iff $A - \{s\}$ is the optimal set of projects.

Proof: 1) infinite capacities on prerequisite edges ensure $A - \{s\}$ is feasible; and 2) the revenue is maximized because the capacity of the cut is minimized:

$$C(A, B) = \sum_{v \in B: p(v) > 0} p(v) + \sum_{v \in A: p(v) < 0} -p(v) = \sum_{v: p(v) > 0} p(v) - \sum_{v \in A} p(v)$$